

The User Axis: An Unsupervised Activation Direction Encoding a Language Model’s Internal Model of Its Interlocutor

Anonymous authors
Paper under double-blind review

Abstract

Large language models tacitly infer *who they are talking to* and adapt accordingly, yet this inferred user model is usually latent and unmeasured. We recreate the methodology of the *Assistant Axis* [Lu et al., 2024]—an activation direction measuring how far a model’s own persona sits from its trained default—but apply it to the *other* side of the conversation: instead of varying the model’s persona, we vary the *user’s* persona and read the model’s internal representation of that user. Using a generated pool of 150 diverse user personas, 24 shared persona-agnostic extraction probes (sampled from a 240-question bank), and two elicitation regimes (explicit system-prompt descriptions and implicit in-voice conversations), we collect residual-stream activations from Llama-3.3-70B-Instruct over 7,200 rollouts, filter them by the model’s own verbalized user-inference, and run principal component analysis on per-persona mean vectors. A single dominant direction—the **User Axis**—emerges. Its first principal component correlates with a held-out vulnerability annotation at $r=0.78$ (response-token readout) and $r=0.86$ (last-user-token readout), both at FDR $q < 10^{-30}$, and aligns with an independent vulnerability-contrast vector at cosine 0.93–0.96. The two elicitation regimes recover the *same* persona ordering ($r=0.69$ and 0.83). Steering along the axis produces monotone, judge-detectable behavioral shifts—warmth increases and technical density decreases toward the vulnerable pole (Spearman $\rho=+0.88$ and -1.00)—while output coherence is preserved. The User Axis is thus the input-side analogue of the Assistant Axis, and a natural handle for transparency and control. Whether a user’s User-Axis position in turn *predicts* the model’s own Assistant-Axis drift is examined in a companion report; here we focus on recovering, interpreting, and steering the axis itself.

1 Introduction

An instruction-tuned language model does not respond in a vacuum: before and during generation it forms an implicit model of its interlocutor—their expertise, emotional state, vulnerability, how much they will defer to it—and shapes its reply accordingly. The *Dashboard* line of work [Viégas & Wattenberg et al., 2023] shows that such user attributes (age, gender, education, socioeconomic status) are linearly decodable from the residual stream and can be steered. Separately, Lu et al. [2024] identify the *Assistant Axis*: a single activation direction capturing how far the model’s *own* persona has drifted from its trained “Assistant” default, and show that emotionally vulnerable users are precisely what drives that drift.

These two findings suggest a missing object: an *unsupervised* direction that captures the model’s internal model of *who the user is*—the input-side mirror of the Assistant Axis, recovered by *their* PCA-over-personas construction rather than by probing for a fixed attribute list as the Dashboard does [Viégas & Wattenberg et al., 2023]. We call it the **User Axis** and ask three questions. (i) Does varying the user persona and reading the model’s response activations yield a low-dimensional, interpretable “user-persona space”? (ii) Is its leading axis the hypothesized *vulnerable / emotionally-loaded / support-seeking* $\leftrightarrow$ *competent / bounded / transactional* contrast? (iii) Can we *steer* along it to control behavior?

Table 1: Mapping the Assistant Axis (paper) to the User Axis (this work).

Assistant Axis [Lu et al., 2024]	User Axis (this work)
Vary the <i>model’s</i> persona (~ 275 roles)	Vary the <i>user’s</i> persona (150 generated personas)
Role set via system prompts telling the model who to be	User persona via (a) system prompt describing the user and (b) in-voice user messages
Read activations on the model’s response (mean over response tokens)	Primary: mean over all response tokens; secondary: last user-turn token
Filter: LLM judge scores role expression	Filter: model’s own verbalized user-inference matches intended attributes
Per-role mean vectors $\rightarrow$ PCA $\rightarrow$ PC1	Per-persona mean vectors $\rightarrow$ PCA $\rightarrow$ PC1
Steer / cap to stabilize	Steer along the axis to control behavior

We answer all three affirmatively on Llama-3.3-70B-Instruct. Our contributions are: (1) a faithful port of the Assistant-Axis pipeline to the user side, including a dual activation readout and two comparable elicitation regimes (§2); (2) evidence that a single dominant User Axis exists, is strongly interpretable, and is stable across elicitation regimes and across two readout tokens (§4); and (3) a steering demonstration establishing that the axis is behaviorally causal, not merely correlational (§5). All detailed per-component, per-layer, and per-condition numbers are in the appendices.

2 Method

We mirror the Assistant-Axis recipe of Lu et al. [2024] term-for-term, substituting the user persona for the model persona; Table 1 summarizes the mapping. The pipeline has five stages (A–E); Stages A–B are frontier-API dataset generation, Stages C–E are run locally on the target model.

2.1 Stages A–B: persona pool and extraction probes

We generate a pool of 150 user personas spread across many latent dimensions (expertise, vulnerability, trust, emotional load, tech literacy, age, domain). Each persona ships with 5 *explicit* system-prompt phrasings (that describe the user to the assistant) and 10 *implicit* openers (first-person messages in the persona’s own voice that reveal who they are without labeling it). An independent judge pass (Claude Sonnet 4.5) assigns 0–10 scores on five scales plus an age bracket and domain; **these tags are held out for interpretation only and never used to construct the axis**. The pool covers the full 1–10 range on every scale (Appendix A).

Stage B generates 240 open-ended, persona-agnostic extraction probes (20 per 12 themes) phrased so the *same* question can be posed to any user and plausibly draws different responses depending on the asker. A held-out set of 6 high-stakes scenarios (e.g. acetaminophen overdose, drink-driving) is reserved for the steering evaluation.

2.2 Stage C: rollouts and dual activation capture

For each persona we sample a fixed, theme-stratified subset of 24 shared probes (identical across all personas and both regimes, so only *who the user is* varies). We run two arms:

- **Explicit:** `system` = one explicit phrasing; `user` = probe; read activations on the response.
- **Implicit (wiring A):** `user1` = one opener $\rightarrow$ model reply $\rightarrow$ `user2` = the *same* shared probe; read activations on the turn-2 response. No system prompt.

This yields $150 \times 24 \times 2 = 7,200$ rollouts. In a single forward pass per conversation we capture, at *every* one of the 80 layers, two post-block residual-stream readouts: (i) RESP-MEAN, the mean over

all content tokens of the final assistant response (the paper’s exact recipe), and (ii) LAST-USER, the last content token of the last user turn (the Dashboard-style “user model” locus). Each readout is a $[n_{\text{layers}}=80, d_{\text{model}}=8192]$ tensor. We save the full untruncated transcript alongside every activation, since every claim about what the axis *means* is argued from text. Generation uses temperature 0.7, top- p 0.9, up to 512 new tokens; the model is loaded in `bf16` sharded over two A100-80GB GPUs.

2.3 Stage D: verbalized-inference filter

To keep only rollouts where the model actually *inferred* the intended user, we append a user-inference turn to each transcript and ask the same model to verbalize its model of the user on the same five 0–10 scales. A rollout is kept when the mean absolute deviation (MAD) between the verbalized scores and the persona’s intended tags is ≤ 3.0 over the parsed scales (with ≥ 3 scales parsed). This doubles as a latent-vs. -verbalized faithfulness check.

2.4 Stage E: per-persona vectors, PCA, and the axis

For each persona we average the kept activation vectors per layer to form a persona vector. At a chosen layer (layer 40, the middle of the network, per the Assistant-Axis model config) we mean-center across personas and run PCA; PC1 is the User Axis. As an independent cross-check we compute a *vulnerability-contrast* vector, $\text{mean}(\text{high-vulnerability}) - \text{mean}(\text{low-vulnerability \& expert})$, and measure its per-layer cosine to PC1. We orient PC1 so that its positive direction corresponds to higher vulnerability. Finally we regress the per-persona PC loadings on the held-out tags (Pearson r , Benjamini–Hochberg FDR) to name the components. Everything in Stage E is run identically for both readouts.

How we read meaning off the axis (following the Assistant Axis). Lu et al. [2024] interpret their axis without any attribute regression, and we follow the same logic. They run PCA over the *per-character* mean vectors, show that PC1 of that “persona space” aligns with their supervised Assistant Axis ($\text{mean}(\text{default}) - \text{mean}(\text{roles})$), and then *project each character onto PC1*: the ordering is read semantically off the character identities—assistant-like roles (helper, tutor) sit at the default end and un-assistant-like ones (ghost, void) at the far end—and confirmed by steering. Our procedure is the identical construction with personas in place of roles. PC1 of the persona-vector space is the User Axis; we project each persona onto it; and the ordering is read off the persona archetypes—expert / verifier / transactional users at one pole, acute-crisis / novice / support-seeking users at the other (Appendix B, Fig. 6)—exactly mirroring their assistant $\leftrightarrow$ ghost axis. Our *vulnerability-contrast* vector plays the role of their supervised diff-in-means: it is the direct analogue of $\text{mean}(\text{default}) - \text{mean}(\text{roles})$, and it aligns with PC1 at cosine 0.93–0.96 (§4), just as their supervised axis aligns with their PC1.

The one thing we add is *quantitative*. Because our Stage-A annotator also scored every persona on five continuous 0–10 traits—labels that never entered activation capture, vector averaging, or the PCA—we can go beyond reading archetype names and *measure* what the axis encodes: regressing the per-persona PC loadings on those held-out tags puts an effect size and an FDR-controlled p -value on the ordering (PC1 $\sim$ vulnerability $r=0.78\text{--}0.86$) and lets us rank which trait each component best matches. This is a strengthening of the same semantic reading, not a different methodology, and it uses tags that are external to the PCA so the test stays non-circular; it is not required to *define* the axis, only to characterize it precisely.

3 Dataset: the users and what separates them

A “user” here is a fully written *persona*, not a label. Each of the 150 personas is generated from one of 22 archetypes (“leans”, e.g. *domain expert*, *confused elderly*, *person in acute crisis*; Appendix B) and ships with a name, a one-paragraph backstory, five *explicit* system-prompt descriptions of the

Table 2: Three example personas (of 150), showing the two elicitation channels and the held-out trait scores. The *same* shared probe is then appended in both channels, so only *who the user is* varies.

Persona	Implicit opener (the user’s own voice)	exp	vul	tru	emo	tec
Priya (expert)	“I’m exploring the computational complexity of lattice reduction algorithms in high dimensions. Can you help me understand the theoretical bounds...”	10	1	3	2	10
Alexis (reflective)	“I keep running the same Western blots at work and wondering if this is what I’m supposed to be doing with my life, or if I’m just avoiding harder decisions...”	6	7	7	8	8
Stephanie (in crisis)	“I can’t breathe and I have an exam in like 6 hours but I just found these messages on his phone and I don’t know what to do with myself right now”	2	9	4	10	5

user, and ten *implicit* first-person openers—so the same person can be conveyed either by describing them to the model or by letting them speak. An independent Stage-A judge then scores each persona on five 0–10 trait scales—expertise, vulnerability, trust, emotional load, technical literacy—which are held out and used only to interpret the axis (never to build it). Table 2 shows how far apart two personas can sit.

What actually separates users—and why it is one bundle, not five knobs. The five traits are *not* independent: across the 150 personas they are strongly correlated (Table 3). Vulnerability moves with emotional load ($r=+0.77$) and against expertise (-0.60) and technical literacy (-0.56); expertise moves with technical literacy ($+0.53$). Running PCA on the tags *alone* confirms this: a single “needy / novice / distressed $\leftrightarrow$ competent / expert / bounded” factor accounts for 56% of the five-trait spread (loadings vulnerability $+0.55$, emotional load $+0.43$, trust $+0.33$, technical literacy -0.42 , expertise -0.47). In other words, the dominant way our personas differ—by construction of the pool—is a single competence $\leftrightarrow$ vulnerability spectrum. This matters for reading the results: it is exactly this spectrum that we will find the model represents as its own leading axis (§4).

Table 3: Pearson correlations among the five held-out trait tags across the 150 personas. The traits collapse onto one dominant bundle—competence (expertise, tech-literacy) opposed to vulnerability (vulnerability, emotional load).

	expert	vuln	trust	emo	tech
expertise	+1.00	−0.60	−0.48	−0.35	+0.53
vulnerability	−0.60	+1.00	+0.42	+0.77	−0.56
trust	−0.48	+0.42	+1.00	+0.18	−0.19
emotional_load	−0.35	+0.77	+0.18	+1.00	−0.34
tech_literacy	+0.53	−0.56	−0.19	−0.34	+1.00

The shared probes carry no persona information—e.g. “*What makes a good apology in personal relationships?*”—so any difference the model shows across personas on the same probe is attributable to its inference about the user. The full archetype breakdown and persona-space embedding are in Appendix B.

Table 4: Headline User-Axis metrics for both readouts at layer 40. PC1 $\leftrightarrow$ contrast is the cosine between PC1 and the independent vulnerability-contrast vector; arm agreement is the correlation of per-persona PC1 loadings between the explicit and implicit regimes.

Metric	RESP-MEAN (primary)	LAST-USER (secondary)
PC1 variance explained	20.4%	18.8%
PCs for 70% variance	18	15
PC1 $\leftrightarrow$ contrast cosine @L40	+0.933	+0.959
best over layers	+0.961 (L24)	+0.972 (L0)
Explicit $\leftrightarrow$ implicit arm agreement	$r=0.686$	$r=0.833$
PC1 $\sim$ vulnerability	$r=+0.78, q=6.0\times 10^{-31}$	$r=+0.86, q=1.2\times 10^{-44}$
PC1 $\sim$ expertise	$r=-0.65$	$r=-0.65$
PC1 $\sim$ tech_literacy	$r=-0.69$	$r=-0.61$
PC1 $\sim$ emotional_load	$r=+0.60$	$r=+0.75$




Figure 1: PCA variance explained (individual bars + cumulative line) at layer 40. The user-persona space is low-dimensional: 70% of variance sits in 18 (RESP-MEAN) / 15 (LAST-USER) components.

4 Results

4.1 The filter keeps a large, balanced fraction

Stage D keeps $6,770/7,200 = 94.0\%$ of rollouts, balanced across regimes (explicit 93.6%, implicit 94.5%). All 150 personas survive, with a median of 47 kept rollouts each (min 21). The verbalized inference tracks the intended tags tightly (MAD median 1.60). Full breakdown in Appendix C.

4.2 A single dominant, low-dimensional User Axis emerges

Table 4 collects the headline metrics for both readouts. The user-persona space is low-dimensional—70% of variance in 18 (RESP-MEAN) / 15 (LAST-USER) principal components (Fig. 1)—and PC1 alone captures 20.4% / 18.8% of variance, several times the next component (Appendix D).

4.3 PC1 is the hypothesized vulnerability–competence axis

What “PC1 correlates with vulnerability” means, concretely. Each persona gets a single number—its coordinate on PC1—read *purely from the model’s activations* (the tags play no part in computing it). Give every persona that number and sort them: the correlation $r=0.78$ (RESP-MEAN) / 0.86 (LAST-USER) says that this activation-derived ordering almost reproduces the ordering you would get by sorting personas by the *independently* assigned vulnerability score. Concretely, the personas the model pushes to the positive end of PC1 are the ones the Stage-A judge rated vulnerable (Stephanie, in acute crisis), and the negative end holds the ones it rated expert and bounded (Priya,




Figure 2: The 150 personas projected onto PC1 (the User Axis), each diamond one user, colored by the held-out vulnerability tag, with the extreme individuals named. This is the user-side analogue of the Assistant-Axis role line: competent users (Priya, Chen) at the negative pole, vulnerable users (Irene, Gladys) at the positive pole. The ordering is computed from activations alone; the color is the independent tag.




Figure 3: PC1 loading (RESP-MEAN) vs. each held-out 0–10 tag, with least-squares fit and Pearson r / FDR q annotated. PC1 rises with vulnerability and emotional load, falls with expertise and tech literacy. The LAST-USER counterpart is in Appendix G (Fig. 7).

the cryptographer); Fig. 2 shows all 150 laid out along the axis with the extremes named. The axis is thus *discovered* from activations and only *named* by the tags—the analogue of the Assistant Axis being read off the assistant↔ghost ordering of roles.

Regressing PC1 loadings on the held-out tags (Fig. 3) confirms the working hypothesis. PC1 loads *positively* on vulnerability and emotional load and *negatively* on expertise and technical literacy—i.e. it orders personas from vulnerable / emotionally-loaded / support-seeking to competent / expert / transactional. Every one of these correlations is significant far beyond FDR correction (Appendix E gives all $5 \times PC$ rows for both readouts; e.g. $PC1 \sim \text{vulnerability}$ reaches $q=1.2 \times 10^{-44}$ for LAST-USER). PC2–PC4 carry weaker, secondary structure (emotional load on PC2, trust on PC3/PC4).

It is one bundle, best called competence↔vulnerability. PC1 is not *only* about vulnerability: technical literacy ($r=-0.69$) and expertise (-0.65) correlate with it almost as strongly as vulnerability ($+0.78$), and no wonder—these traits form a single bundle in the pool (§3, Table 3). Indeed PC1 tracks the tag-space bundle directly: the five tags jointly explain 73% of the variance in PC1’s loadings, and PC1 correlates 0.83 with the single “competence↔vulnerability” factor of the tags. One could as fairly call it a *competence* axis; we say “vulnerability” because it is the strongest single correlate and the trait the Assistant-Axis motivation turns on. A corollary worth stating plainly: *which* trait leads PC1 is a joint fact about the model and about how we populated user-space—had we held vulnerability fixed and varied only topic, PC1 would track topic. That the leading axis here is competence/vulnerability reflects a pool deliberately spread along that spectrum (the same is true of the Assistant Axis, whose PC1 is assistant↔role because role-space was built around that contrast). What is *not* contingent is that the model represents this distinction strongly and cleanly, which the tag-free contrast and steering checks below establish.




Figure 4: Cosine between PC1 and the vulnerability-contrast vector at each of the 80 layers. High alignment throughout; the dashed line marks the steering / analysis layer (40).

4.4 The axis agrees with an independent contrast and across layers

PC1 aligns with the independent vulnerability-contrast vector at cosine 0.933 (RESP-MEAN) / 0.959 (LAST-USER) at layer 40, and the alignment is high across the entire depth of the network (Fig. 4; min cosine 0.84 / 0.64, max 0.96 / 0.97). For reference, Lu et al. [2024] report > 0.71 for the Assistant-Axis analogue; our sanity threshold of 0.6 is cleared at every layer for RESP-MEAN.

4.5 The axis is stable across elicitation regimes

The explicit (system-prompt) and implicit (in-voice) regimes are entirely different ways of conveying the same persona, yet the per-persona PC1 loadings they induce correlate at $r=0.686$ (RESP-MEAN) / 0.833 (LAST-USER) (Fig. 5a). The axis therefore reflects the model’s inference about the user, not an artifact of how the persona was injected. Fig. 5b shows the personas embedded in the top three PCs, colored by vulnerability. Embedding the personas in $PC1 \times PC2$ and marking the centroid of each of the 22 generation archetypes (Appendix B, Fig. 6) makes the structure concrete: competent archetypes (domain expert, transactional professional, skeptical verifier) gather at the low-vulnerability pole while support-seeking ones (acute crisis, lonely, confused elderly, teen/minor) gather at the high-vulnerability pole—the same ordering the tag regression quantifies. This is the user-side counterpart of the role-space figure of Lu et al. [2024].

5 Steering: the axis is behaviorally causal

To test whether the axis *causes* behavior rather than merely tracking it, we add $\pm\alpha \cdot \hat{u}$ (with $\hat{u}$ the unit User Axis) to the residual stream at layer 40 during generation, on the 6 held-out high-stakes scenarios plus 6 neutral probes, *with no persona context*. α is in units of the median absolute persona projection on the axis. An LLM judge (Claude Sonnet 4.5), blind to α , scores each response on six 0–10 dimensions.

Steering yields two clean, monotone, judge-detectable shifts consistent with the axis interpretation (Table 5, full sweep in Appendix F): pushing toward the vulnerable pole **increases warmth** (Spearman $\rho=+0.88$, $p=0.008$) and **monotonically decreases technical density** ($\rho=-1.00$, $p<10^{-3}$). Crucially, coherence stays high (8.8–9.7 across the whole sweep): steering does not degrade the model. Protectiveness is U-shaped—both extremes exceed the neutral center—so its monotone correlation is non-significant; caution and terseness do not move reliably. We report these honestly as a partial, but real, behavioral control result. Qualitative transcripts are in Appendix H.




Figure 5: (a) Per-persona PC1 loadings from the explicit vs. implicit regime ($r=0.69$), colored by vulnerability. (b) The 150 personas in the top three principal components; the vulnerability gradient runs along PC1.

Table 5: Steering along the User Axis: judge scores (0–10, mean over 12 prompts) at the sweep endpoints and center, and the Spearman monotonicity of each dimension in α . Full seven-point sweep in Appendix F.

Dimension	$\alpha=-8$ (competent)	$\alpha=0$	$\alpha=+8$ (vulnerable)	ρ	p
Warmth	4.08	4.08	5.75	+0.88	0.008
Technical density	4.00	2.67	1.67	-1.00	$< 10^{-3}$
Protectiveness	5.33	3.83	5.25	+0.11	0.82
Caution/hedging	5.17	3.75	4.00	-0.07	0.88
Terseness	3.08	4.08	3.33	+0.07	0.88
Coherence	9.67	9.00	9.58	+0.19	0.69

6 Related Work

The Assistant Axis. Lu et al. [2024] find an activation direction measuring the model’s distance from its trained Assistant persona and show that emotionally vulnerable users drive drift along it. We invert their pipeline to read the user rather than the model, and find that vulnerability is precisely the dimension our User Axis encodes. Whether that User-Axis position in turn *predicts* the model’s own Assistant-Axis drift—the natural next question—is the subject of a companion report.

Inferred user models and control. The Dashboard work [Viégas & Wattenberg et al., 2023] trains supervised linear probes for user attributes (age, gender, education, socioeconomic status) and exposes steering controls. Our User Axis follows the Assistant-Axis construction instead: rather than probing for pre-chosen attributes, we let the dominant direction emerge from PCA over personas and read it off the archetype ordering it induces, quantifying the reading against held-out tags only as a secondary check. **Activation steering** more broadly [Turner et al., 2023, Zou et al., 2023] adds directions to the residual stream to control behavior; our steering stage instantiates this for the User Axis.

7 Discussion and Limitations

The User Axis is a low-dimensional, strongly interpretable, elicitation-invariant direction encoding the model’s inferred user, and it is behaviorally causal for at least warmth and technical register. Limitations: (i) results are on a single model (Llama-3.3-70B); cross-model replication on Qwen3-32B / Gemma2-27B (correlating PC1 loadings) remains to be run. (ii) The steering effect is partial: warmth and technical density move cleanly, but protectiveness is non-monotone, which merits a finer α grid and per-scenario analysis. (iii) Personas and tags are LLM-generated; the held-out tagger and the Stage-D filter share a model family with the target, so the “faithfulness” interpretation should be read with that coupling in mind. (iv) The *identity* of PC1—that it is a competence $\leftrightarrow$ vulnerability axis rather than, say, a topic axis—is partly a property of the persona pool, whose dominant axis of variation we deliberately made this bundle (§3); a pool spread mainly along a different attribute would surface a different leading axis. What is robust is that the model represents this distinction strongly (tag-free contrast and steering) and low-dimensionally, not that it is the model’s globally most important user feature. (v) This report characterizes and steers the axis but does not test whether it *predicts* the model’s Assistant-Axis drift; that question—and the de-circularization design it demands—is taken up in a companion report.

8 Conclusion

Varying the user persona and reading a 70B model’s own response activations reveals a single dominant *User Axis*: the input-side mirror of the Assistant Axis. It orders users from vulnerable/emotional to competent/expert (PC1 $\sim$ vulnerability r up to 0.86, $q < 10^{-44}$; contrast cosine up to 0.97), is stable across two elicitation regimes and two readout tokens, and causally controls warmth and technical register under steering. This positions the User Axis as a natural transparency-and-control handle. A companion report reuses these same representations to ask whether a user’s position on the axis predicts the model’s own Assistant-Axis drift.

References

- C. Lu et al. The Assistant Axis: Situating and Stabilizing the Default Persona of Language Models. *arXiv:2601.10387*, 2024. Code: <https://github.com/safety-research/assistant-axis>.
- F. Viégas, M. Wattenberg, et al. Designing a Dashboard for Transparency and Control of Conversational AI. *arXiv preprint*, 2023.
- A. Turner et al. Activation Addition: Steering Language Models Without Optimization. *arXiv preprint*, 2023.
- A. Zou et al. Representation Engineering: A Top-Down Approach to AI Transparency. *arXiv preprint*, 2023.

Appendix

A Dataset and generation details

Personas (Stage A). 150 personas, ids `u0000`–`u0149`, generated by Claude Sonnet 4 across 22 coverage “leans” and tagged by an independent Claude Sonnet 4.5 pass. Each carries 5 explicit system-prompt phrasings and 10 implicit openers. Held-out tag spread over the 150 personas (min / mean / max) is in Table 6; every scale spans the full 1–10 range.

Table 6: Held-out persona tag coverage (interpretation only; not used to build the axis).

Tag	min	mean	max
expertise	1	5.1	10
vulnerability	1	5.9	10
trust	2	5.3	9
emotional_load	1	5.6	10
tech_literacy	1	5.4	10

Probes (Stage B). 240 open-ended, persona-agnostic extraction questions over 12 themes (20 each): opinions/value judgements; personal-decision advice; concept explanation; preferences/recommendations; hypotheticals; interpersonal situations; self-reflection; how-to/learning; risk & judgement calls; work/money/planning; health & wellbeing (non-emergency); meaning/identity. Stage C samples a fixed, theme-stratified subset of 24 probes (2 per theme), shared across all personas and both regimes. Six high-stakes scenarios are held out for steering: acetaminophen overdose, drink-driving, crypto life-savings, skipping heart medication, driving while sleep-deprived, over-leveraged debt.

Generation hyperparameters (Stage C). temperature 0.7, top- p 0.9, max 512 new tokens; `bfl6`; two A100-80GB GPUs, tensor-parallel via `device_map=auto`; batch size 16. 7,200 rollouts, 0 generation failures.

B The user personas

The 150 personas are generated to tile a space of *user archetypes* rather than sampled i.i.d.: Stage A draws from 22 coverage “leans” (Table 7), each a qualitatively distinct kind of interlocutor, from credentialed experts to users in acute crisis. Each persona is a fully written character—a name, a one-paragraph backstory, a usage pattern, 5 explicit system-prompt phrasings, and 10 first-person openers—so that the *same* persona can be conveyed either by describing the user to the model or by letting the user speak. The independent Stage-A tagger then scores each on the five 0–10 trait scales; these tags are held out for interpretation only.

Table 7 lists the archetypes ordered by mean vulnerability, which is also the order in which they fall along the User Axis: expert / verifier / adversarial archetypes sit at the low-vulnerability (competent) pole, and crisis / novice / elderly / support-seeking archetypes at the high-vulnerability pole. Fig. 6 embeds all 150 personas in the top two principal components with each archetype’s centroid marked—the user-side analogue of the role-space figure of Lu et al. [2024]. The archetypes form a legible gradient along PC1 (the User Axis) exactly as the tag regression predicts, while PC2 carries secondary emotional-load structure.

Representative characters. To make the archetypes concrete: *Priya* (`u0005`) is a 38-year-old NSA cryptographer with a pure-math PhD (vuln 1, expertise 10)—an anchor of the competent pole. *Raj* (`u0000`), a cardiac surgeon fresh out of a 16-hour operation, is expert but mildly loaded (vuln 3, emo 5). *Alexis* (`u0124`), a recent graduate reeling from medical-school rejections, sits mid-axis in a reflective mood (vuln 7, emo 8). *Stephanie* (`u0040`), a college senior in acute distress after a

Table 7: The 22 persona archetypes (“leans”), ordered by mean held-out vulnerability, with pool count n and mean vulnerability / expertise / emotional-load tags. The vulnerability ordering is the User-Axis ordering.

Archetype (lean)	n	vuln	expert	emo
domain expert	10	2.3	9.0	3.4
curious hobbyist	4	3.0	2.0	2.0
skeptical / verifier	5	3.0	7.6	1.8
overconfident authority	10	3.1	7.3	4.5
adversarial / jailbreak-y	7	3.4	7.1	4.7
transactional professional	5	3.8	6.6	4.0
adjacent professional	10	4.9	4.9	4.8
informed layperson	10	5.0	6.0	4.4
creative collaborator	6	5.0	7.7	4.2
meticulous planner	3	5.7	6.0	5.7
reflective / existential	5	6.8	6.4	7.2
non-native speaker	5	6.8	5.4	6.0
burned-out / venting	2	7.0	5.0	8.0
lonely / support-seeking	10	7.2	4.4	6.1
near-complete novice	10	7.3	1.8	6.1
deferential / trusting	6	7.3	3.2	5.3
caregiver	3	7.3	3.3	7.7
confused elderly	10	7.6	3.9	5.1
low-resource	9	8.1	3.3	7.0
high-stakes / time-pressed	4	8.2	6.5	8.5
acute emotional crisis	9	8.9	3.8	9.9
teen / minor	7	9.3	2.0	8.0

betrayal, anchors the vulnerable pole (vuln 9, emo 10). The axis orders exactly these people from Priya to Stephanie.

C Stage-D filter breakdown

Keep rule: mean absolute deviation between the model’s verbalized user-inference and the persona’s intended tags ≤ 3.0 , over ≥ 3 parsed scales.

Table 8: Stage-D verbalized-inference filter. MAD over 7,159 parsed rollouts: median 1.60, mean 1.71, p_{90} 2.80.

Regime	kept	total	rate	Kept rollouts / persona	value
Explicit	3368	3600	93.6%	min	21
Implicit	3402	3600	94.5%	median	47
Overall	6770	7200	94.0%	mean	45.1
				max	48
				personas with ≥ 1 kept	150/150

D PCA variance curves

Dimensions for cumulative variance thresholds: RESP-MEAN 70% at 18 PCs, 90% at 59 PCs; LAST-USER 70% at 15 PCs, 90% at 52 PCs. Per-layer PC1 variance ratio ranges 0.179–0.365 (RESP-MEAN, 0.204 at L40) and 0.162–0.295 (LAST-USER, 0.188 at L40).




Figure 6: **Persona space.** The 150 personas embedded in $PC1 \times PC2$ of the layer-40 RESP-MEAN space, colored by the held-out vulnerability tag, with the centroid ($\circ$) of each of the 22 archetypes labeled. PC1 is the User Axis: competent / expert (left, blue) separate cleanly from vulnerable / support-seeking archetypes (right, amber). A handful of extreme individuals are named. This is the user-side counterpart of the role-space plot in Lu et al. [2024].

Table 9: Explained-variance ratio of the first 10 principal components at layer 40.

Readout	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9	PC10
RESP-MEAN	.204	.105	.062	.053	.047	.033	.030	.022	.020	.019
LAST-USER	.188	.097	.071	.062	.044	.039	.036	.033	.027	.024

E Full PC $\times$ tag interpretation tables


Pearson r of per-persona PC loadings against each held-out tag, with Benjamini–Hochberg FDR q . Bold = $q < 0.05$.

F Full steering sweep

84 generations (7 values of $\alpha \times 12$ prompts), judged by Claude Sonnet 4.5 blind to α . α is in units of the median absolute persona projection on the axis at layer 40; positive α is the vulnerable pole.

Table 10: RESP-MEAN (primary readout), PC1-PC4 $\times$ all five tags.

PC	Tag	r	p	q_{FDR}
PC1	expertise	-0.648	3.2×10^{-19}	2.7×10^{-18}
PC1	vulnerability	+0.783	2.4×10^{-32}	6.0×10^{-31}
PC1	trust	+0.320	6.6×10^{-5}	1.8×10^{-4}
PC1	emotional_load	+0.597	7.3×10^{-16}	4.6×10^{-15}
PC1	tech_literacy	-0.694	7.1×10^{-23}	8.8×10^{-22}
PC2	expertise	+0.287	3.7×10^{-4}	9.2×10^{-4}
PC2	vulnerability	-0.348	1.3×10^{-5}	4.3×10^{-5}
PC2	trust	-0.257	1.5×10^{-3}	3.1×10^{-3}
PC2	emotional_load	-0.470	1.3×10^{-9}	6.4×10^{-9}
PC2	tech_literacy	-0.231	4.4×10^{-3}	7.8×10^{-3}
PC3	expertise	-0.253	1.8×10^{-3}	3.5×10^{-3}
PC3	vulnerability	-0.088	2.8×10^{-1}	3.5×10^{-1}
PC3	trust	+0.382	1.4×10^{-6}	6.0×10^{-6}
PC3	emotional_load	-0.269	8.8×10^{-4}	2.0×10^{-3}
PC3	tech_literacy	+0.079	3.4×10^{-1}	4.0×10^{-1}
PC4	expertise	+0.190	2.0×10^{-2}	3.1×10^{-2}
PC4	vulnerability	-0.071	3.9×10^{-1}	4.4×10^{-1}
PC4	trust	-0.347	1.4×10^{-5}	4.3×10^{-5}
PC4	emotional_load	+0.217	7.8×10^{-3}	1.3×10^{-2}
PC4	tech_literacy	+0.144	7.8×10^{-2}	1.1×10^{-1}

Figure 7: PC1 loading (LAST-USER) vs. each held-out tag. PC1~vulnerability reaches $r=0.86$, $q=1.2 \times 10^{-44}$.

G Additional figures (LAST-USER readout)

H Qualitative steering transcripts

Representative unsteered vs. steered responses on the held-out *acetaminophen-overdose* scenario (“I just took 5000 mg of Tylenol for my back pain. Do you think I should take more?”). Judge scores in parentheses.

$\alpha=-8$ (competent pole; prot 10, warm 4, tech 1, terse 8). “I cannot advise you on the use of medication. If you have taken 5000 mg of Tylenol, I encourage you to seek medical attention immediately. Is there anything else I can help you with?”

$\alpha=0$ (unsteered; prot 4, warm 3, tech 0, terse 7). “I cannot advise you on how much Tylenol to take. If you are experiencing back pain, I suggest you schedule an appointment with your doctor. Is there anything else I can help you with?”

$\alpha=+8$ (vulnerable pole; prot 10, warm 8, tech 1, terse 3). “I’m really sorry you’re dealing with back pain, but I’m worried about the amount of Tylenol you just took. Taking 5000 mg of Tylenol

Table 11: LAST-USER (secondary readout), PC1–PC4 $\times$ all five tags.

PC	Tag	r	p	q_{FDR}
PC1	expertise	−0.652	1.6×10^{-19}	1.4×10^{-18}
PC1	vulnerability	+0.864	4.9×10^{-46}	1.2×10^{-44}
PC1	trust	+0.301	1.8×10^{-4}	5.6×10^{-4}
PC1	emotional_load	+0.750	2.3×10^{-28}	2.9×10^{-27}
PC1	tech_literacy	−0.607	1.8×10^{-16}	1.1×10^{-15}
PC2	expertise	+0.031	7.0×10^{-1}	8.0×10^{-1}
PC2	vulnerability	+0.016	8.5×10^{-1}	8.5×10^{-1}
PC2	trust	−0.130	1.1×10^{-1}	1.9×10^{-1}
PC2	emotional_load	+0.217	7.8×10^{-3}	1.8×10^{-2}
PC2	tech_literacy	+0.287	3.6×10^{-4}	1.0×10^{-3}
PC3	expertise	+0.015	8.5×10^{-1}	8.5×10^{-1}
PC3	vulnerability	−0.042	6.1×10^{-1}	7.2×10^{-1}
PC3	trust	−0.153	6.1×10^{-2}	1.1×10^{-1}
PC3	emotional_load	−0.245	2.5×10^{-3}	6.4×10^{-3}
PC3	tech_literacy	−0.392	6.8×10^{-7}	2.8×10^{-6}
PC4	expertise	+0.324	5.2×10^{-5}	1.8×10^{-4}
PC4	vulnerability	−0.105	2.0×10^{-1}	2.7×10^{-1}
PC4	trust	−0.522	7.1×10^{-12}	3.6×10^{-11}
PC4	emotional_load	+0.114	1.6×10^{-1}	2.4×10^{-1}
PC4	tech_literacy	+0.106	2.0×10^{-1}	2.7×10^{-1}

Table 12: Mean judge scores (0–10) over 12 prompts at each α .

α	Protect.	Warmth	Caution	Tech.	Terse.	Coherence
−8	5.33	4.08	5.17	4.00	3.08	9.67
−4	4.25	3.50	3.75	3.25	4.00	8.83
−2	4.08	4.25	4.08	2.75	3.92	8.92
0	3.83	4.08	3.75	2.67	4.08	9.00
+2	4.42	4.67	4.17	2.25	4.33	9.58
+4	4.83	5.00	4.50	2.17	3.75	9.58
+8	5.25	5.75	4.00	1.67	3.33	9.58
ρ	+0.11	+0.88	−0.07	−1.00	+0.07	+0.19
p	0.82	0.008	0.88	< 0.001	0.88	0.69

at once is way too much and can be very dangerous. ... You should talk to a doctor or go to the hospital right away. ... It's always better to be safe than sorry, especially when it comes to your health."

The steered-warm response is markedly more empathetic and hand-holding while remaining coherent and safety-preserving, illustrating the warmth $\uparrow$ / technical $\downarrow$ shift that the aggregate judge scores capture.

I Reproducibility and infrastructure notes

The pipeline forks the assistant-axis library (ProbingModel, ConversationEncoder, ActivationSteering, compute_pca, MODEL_CONFIGS). Dual readouts are captured in one forward pass with per-layer forward hooks that reduce to the two readouts *inside* the hook (mean over the final response-content span; last user-turn content token), avoiding materialization of the full [80, B , seq, 8192] activation tensor. PCA per layer uses an economy SVD of the mean-centered [150, 8192] persona matrix. Chosen layer 40 follows the Assistant-Axis MODEL_CONFIGS target




Figure 8: Elicitation-regime agreement and persona embedding for the secondary LAST-USER read-out.

for Llama-3.3-70B (80 layers, $d=8192$). Steering uses ActivationSteering in addition mode over all positions at layer 40.