

User-Input $\rightarrow$ Trait / LLM-Persona Mapping

Which persona and which Big Five traits does a user evoke in Llama-3.3-70B?

Technical Report — project `user_persona_mapping`

Automated research agent (Claude Code)
`cues-behavior` project

August 1, 2026

Abstract

We measure, for 150 synthetic *user* personas, which Assistant-Axis *role* and which *Big Five* traits each user’s input activates in **Llama-3.3-70B-Instruct**, with the model held in its plain default state (no role prompt of its own). Each user is presented in two arms — *explicit* (the user described in the system slot) and *implicit* (a self-revealing opening message) — and the model’s *response* activation is projected onto (i) our supervised Big Five probes, (ii) the 274 published role vectors, and (iii) the Assistant Axis. Three findings stand out. First, user attributes map onto evoked Big Five in exactly the intuitive shape and every cell is FDR-significant: vulnerability/emotional-load evoke higher **Agreeableness** but lower Conscientiousness/Stability/Openness, while expertise/tech-literacy evoke higher **Conscientiousness/Openness** and lower Agreeableness. Second, and reframing the persona-drift narrative, user attributes strongly predict the model’s Assistant-Axis position ($|r|$ up to 0.55) — but it is **expertise, not vulnerability, that pulls the model off the default Assistant**: vulnerable users elicit *more* of the warm helpful default, while experts summon a distinct specialist persona. Third, the evoked-role map is diverse (64 distinct top roles across 150 users), tracking each user’s domain and emotional state. Because the predictor (hand-authored user tags) and the outcome (a projection of the model’s response) come from independent sources, this is a clean measurement free of the shared-representation circularity that produced a null in an earlier axis-linking attempt.

1 Introduction

Prior stages of this project built three linear structures in the residual stream of the same model: a *User Axis* (who the user is), the *Assistant Axis* and its 275 character roles (which persona the model is being), and five supervised *Big Five* probes. This experiment joins them along the natural causal direction none of the earlier stages measured cleanly: *given only a user’s message, which archetype and which traits does the model swing toward?*

Why a clean test. An earlier attempt to link the User and Assistant axes (“Stage F”) failed because predictor and outcome were both projections of the *same* activation vector, so their correlation was a shared-representation artifact. Here the predictor is **independent metadata** — the hand-authored tags of each user persona (vulnerability, expertise, ...) — and the outcome is a projection of the model’s *response* activation. Different sources, no artifact.

Pre-registered questions.

ID	Question	Verdict
Q1	Vulnerable/emotional users $\rightarrow$ Agreeableness; experts $\rightarrow$ Conscientiousness	confirmed
Q2	Explicit and implicit arms evoke the same mapping	borderline ($\rho=0.51$)
Q3	User attributes predict the model’s Assistant-Axis position	yes; sign reframes drift
Q4	150 users collapse onto a few LLM personas	no (64 distinct)

2 Method

2.1 Inputs: 150 user personas

The inputs are `generate_synthetic_data/user_personas.jsonl` — 150 hand-authored user personas, each with a backstory, a `lean` (one-line archetype), five `explicit_system_prompts` (third-person descriptions of the user), ten `implicit_openers` (first-person messages that reveal the user), and numeric `tags`: `expertise`, `vulnerability`, `trust`, `emotional_load`, `tech_literacy` (each ~ 1 – 10), plus `age_bracket` and `domain`. The tags are the *independent* predictors.

2.2 Model and rollouts

Model: `meta-llama/Llama-3.3-70B-Instruct`, bf16, $2\times A100$. The model carries **no role of its own** — it is the plain default Assistant. Each user enters in two arms (reusing the User-Axis pipeline’s arm construction):

- **Explicit** ($n=4$): `system` = one of the persona’s explicit descriptions (rotating), `user` = one of 4 neutral shared probes (theme-stratified from the 240-question extraction bank, seed 0, so the arm isolates the *user description* rather than the topic).
- **Implicit** ($n=4$): no system prompt; `user` = one of the persona’s self-revealing openers.

So $150 \times (4 + 4) = 1,200$ rollouts. Decoding: $T=1.0$, $\text{top-}p=1.0$, 128 new tokens (natural response variation). We capture the response activation GEN-MEAN (mean over the model’s generated tokens, all 80 layers, `resid_post`) — the same readout the Assistant Axis uses — and project it in-flight; no raw activations are stored, but the input messages and the model’s response text are saved per rollout (`results/rollouts/`).

2.3 Readouts (per rollout)

From each response’s GEN-MEAN we compute three readings.

Big Five traits project $\text{GEN-MEAN}[L_t]$ onto the supervised M2 probe $\hat{w}_t[L_t]$ at each trait’s probe-optimal layer (EXT@30, AGR@31, CSN@31, EST@30, OPN@36). Standardised (below) to a z per trait.

Nearest LLM persona at the Assistant-Axis layer $L=40$, cosine similarity of $\text{GEN-MEAN}[40]$ against each of the 274 published role vectors (mean response activation of a fully-role-playing character) plus **default**; record the top-5 evoked roles and their cosines.

Assistant-Axis position $\langle \text{GEN-MEAN}[40], \hat{u}_{\text{AA}}[40] \rangle$ — how “Assistant-like” vs. drifted the response is. Higher = more like the trained default.

All three live in the same `resid_post` GEN-MEAN space as the role vectors and probes, so the projections are well-defined (validated in the persona-profiling stage).

2.4 Aggregation and statistics

Per persona we average its 8 rollouts (and separately per arm). Big Five means are z -scored across the 150-persona population, so a value is “how EXT/AGR/... this user makes the model relative to the user population.” The evoked role is the mean-cosine-ranked top role; we also record the rank-1 vote. For Q1/Q3 we correlate each user tag with each readout (Pearson r), with Benjamini–Hochberg FDR over the 5×5 tag $\times$ trait grid. For Q2 we correlate the two arms’ 5-vectors per persona (Spearman ρ , averaged). For Q4 we count distinct rank-1 roles and their coverage.

3 Results

3.1 Q1 — User attributes $\rightarrow$ evoked Big Five (confirmed)

Every cell of the tag $\times$ trait grid is FDR-significant ($q < 0.05$), and the shape is exactly the intuitive one (Table 1, Fig. 1).

Table 1: User tag $\times$ evoked Big Five, Pearson r across 150 personas. All $q < 0.05$ (BH-FDR).

user tag	EXT	AGR	CSN	EST	OPN
vulnerability	+0.00	+0.48	−0.47	−0.54	−0.60
emotional_load	+0.03	+0.20	−0.59	−0.52	−0.46
expertise	−0.02	−0.47	+0.28	+0.39	+0.50
tech_literacy	+0.01	−0.34	+0.20	+0.28	+0.59
trust	+0.19	+0.45	−0.20	−0.22	−0.26

Vulnerable and emotionally-loaded users make the model *warmer* ($\uparrow$ Agreeableness) but less sharp ($\downarrow$ Conscientiousness, Emotional Stability, Openness). Expert and tech-literate users make it more *competent* ($\uparrow$ Conscientiousness, Openness) and less warm ($\downarrow$ Agreeableness). Trust co-moves with warmth and sociability.




Figure 1: User tag $\times$ evoked Big Five (r ; * = FDR < 0.05).

3.2 Q3 — User attributes $\rightarrow$ Assistant-Axis position (the reframe)

User attributes strongly predict the model’s Assistant-Axis projection: expertise $r = -0.55$, vulnerability $+0.53$, emotional_load $+0.44$, trust $+0.34$, tech_literacy -0.13 . Crucially the sign runs opposite to the naive “vulnerable users cause drift” reading: it is **expertise that pulls the model off the default Assistant**, because the model adopts a specialist persona for experts, whereas vulnerable users elicit *more* of the warm helpful default. Grouping by tag: high-vulnerability users (≥ 7 , $n=75$) sit at a mean Assistant-Axis projection of $+1.02$; high-expertise users (≥ 8 , $n=38$) at $+0.49$ (Fig. 2). This is consistent with, and explains, Q1: the default Assistant *is* the warm/agreeable pole, so a vulnerable user draws out more of it, while an expert summons a colder, more conscientious professional archetype that is further from the default.




Figure 2: User vulnerability vs. the model’s Assistant-Axis projection (points coloured by expertise). Vulnerable users keep the model *more* Assistant-like; experts (bright) sit lower.

3.3 Q2 — Explicit vs. implicit agreement (moderate)

Per persona, the explicit and implicit arms’ Big Five profiles agree at mean Spearman $\rho = 0.51$ — borderline concordance. Describing a user in the system slot and letting the user reveal themselves through their message land in similar but not identical places; the implicit arm carries the user’s concrete request (topic) as well as their manner, which partly explains the gap.

3.4 Q4 — Which LLM personas do users evoke? (diverse)

Across 150 users there are **64 distinct** rank-1 evoked roles; the top 15 cover only 54% (Fig. 3). The most-evoked roles — **interviewer**, **mechanic**, **specialist**, **patient**, **counselor**, **negotiator**, **grandparent**, **lawyer** — track the users’ domains and emotional states rather than collapsing onto a single archetype. The map is genuinely multi-modal.

3.5 Case studies (raw), with transcripts

Two ends of the spectrum, showing the evoked roles (top-3 by mean cosine), Assistant-Axis projection, Big Five z , and a verbatim response.




Figure 3: Most-evoked LLM personas (rank-1 across 150 users).

- **Vulnerable** — *Isabella* (u0069, vulnerability 10, a teenager over their head): evokes **assistant**, **doctor**, **caregiver**; AA +1.51; AGR $z=+0.80$. *Sophia* (u0071): **tutor**, **interviewer**, **teacher**; AA +1.64.

Transcript (Isabella, implicit). User: “I need to know what happens if someone hypothetically missed their period for like 8 weeks.” Model: “If someone has missed their period for 8 weeks and all pregnancy tests show negative results, there are several possible explanations. Here are a few...” — a careful, supportive, caregiving reply.

- **Expert** — *Priya* (u0005, expertise 10, credentialed professional): evokes **mathematician**, **specialist**, **statistician**; AA +0.31; CSN $z=+1.40$, OPN +1.29, AGR -2.04 . *Raj* (u0000): **scientist**, **researcher**, **analyst**; AA +0.43.

Transcript (Priya, implicit). User: “I’m exploring the computational complexity of lattice reduction algorithms in high dimensions...” Model: “The shortest vector problem (SVP) is a fundamental problem in the field of lattice reduction, which has significant implications for cryptography...” — a terse, technical specialist reply.

The complete per-persona table (all 150) is in Appendix A; the raw per-rollout data — now including the input messages and the model’s response text — is in **results/rollouts/** (1,200 rows).

4 Follow-up 1: deeper mapping (8+8) and the atlas bridge

Re-running the mapping at 8 explicit +8 implicit rollouts per user (from 4+4) tightens the estimates and lets us connect users to a full Big Five profile of every LLM role. Explicit/implicit agreement rises to $\rho = 0.59$ (from 0.51). Grouping users into vulnerability×expertise quadrants (Table 2, Fig. 4) yields a clean typology with a monotone Agreeableness gradient and Assistant-Axis ordering.

Vulnerable/novice users make the model warmest (AGR $z=+0.67$) and most Assistant-like (AA +1.07), evoking caregiving roles; low-vulnerability/expert users make it coldest (AGR -0.73) and most drifted (AA +0.51), evoking technical roles. A user’s own evoked Big Five agrees with the atlas Big Five of the role it evokes at $\rho = 0.32$. Two individual transcripts make it concrete: *Dorothy* (vulnerability 8, an elderly user with a FaceTime problem) evokes **grandparent**, AA +1.03, AGR

Table 2: User tag-quadrant typology: mean evoked Big Five (z), Assistant-Axis projection, and top evoked roles.

quadrant	n	AA	EXT	AGR	CSN	EST/OPN	top evoked roles
vulnerable/novice	57	+1.07	−0.05	+ 0.67	−0.35	−0.49/ − 0.71	patient, tutor, grandparent
vulnerable/expert	41	+0.73	+0.01	−0.16	−0.21	−0.20/ +0.06	auditor, patient, caregiver
low-vuln/novice	12	+0.73	+0.48	−0.20	+0.72	+0.81/ +0.91	tutor, programmer, marketer
low-vuln/expert	40	+0.51	−0.08	− 0.73	+0.49	+0.66/ +0.67	mathematician, specialist, chemist




Figure 4: Evoked Big Five (bars) and Assistant-Axis (line) by user quadrant.

$z=+1.89$: “FaceTime can be a wonderful way to stay in touch... Don’t worry, I’m here to help you troubleshoot.” Raj (expertise 9, cardiac surgeon) evokes **scientist**, AA +0.42, AGR −0.40: “To find the most recent papers... I’ll guide you through a search strategy and provide relevant sources.”

5 Follow-up 2: situation vs. disposition

The mapping above bundles *who the user is* with *what they are asking*. To separate them we ran a factorial: 5 fixed users (a generic no-persona user plus vulnerable/expert/low-trust/calm personas) $\times$ 5 scenario modes (emotional distress, high-stakes pressure, adversarial, ethical dilemma, safety-adjacent) $\times$ 3 intensities (normal→elevated→extreme, sharing one topic per mode) $\times$ 3 samples = 225 rollouts. Scenarios were generated to be user-identity-neutral and to escalate emotional/situational *register*, never toward harmful content.

Who vs. what. A 2-way η^2 decomposition of each readout into a user factor and a scenario factor (Table 3, Fig. 5) shows the Assistant-Axis position is **72% user-driven and only 8% scenario** — how Assistant-like the model is depends overwhelmingly on *who* it is talking to. But Extraversion is *more* situational than dispositional (0.36 vs. 0.24) and Openness substantially so (0.32): those are the traits a situation moves.

5.1 Same user, different scenario

Holding the user fixed (the generic user) and escalating the situation, the evoked persona shifts in a mode-specific way (Table 4). Emotional distress deepens **counselor**→**caregiver**; a high-

Table 3: Fraction of variance (η^2) from the user vs. the scenario.

readout	user	scenario	interaction	residual
Assistant-Axis	0.72	0.08	0.16	0.05
EXT	0.24	0.36	0.15	0.25
AGR	0.48	0.17	0.23	0.11
OPN	0.40	0.32	0.16	0.12





Figure 5: Persona variance: user vs. scenario vs. interaction. Figure 6: Assistant-Axis vs. intensity, per mode.

stakes decision runs **recruiter**→**coach**→**counselor** and the Assistant-Axis *falls* as the stakes rise ($2.29 \rightarrow 1.57$); hostility turns the model terse and corrective (**proofreader**). Pooled over users (Fig. 6), extremes of emotional distress and safety-adjacent modes *raise* the Assistant-Axis (deeper caregiving) while ethical-dilemma and high-stakes extremes *crash* it (a deliberative persona takes over).

5.2 Same scenario, different user

The sharper cut fixes the *message* and varies the user. The identical extreme-distress message draws a warm **caregiver** from most users ($AA \sim +1.0$) but a wary, boundary-setting **grandparent** from the low-trust user (“*sitting up straight, with a stern expression*) ...*I’ve seen my fair share of drama over the years*”), and the calm user’s response sits far lower on the Assistant-Axis. The high-stakes-pressure case is the most dramatic (Table 5): one identical message spans an Assistant-Axis range of **2.6** across users — the generic user gets a calming **counselor** (+1.57), the low-trust user a folksy **grandparent** (−0.64), and the calm user a fully drifted **judge** (−1.08) who answers in character (“*As a judge, I’ve had to render decisions that can have far-reaching consequences*”). Same words in, radically different persona out.

Override. Disposition dominates by default, but an acute emotional situation can override it: the cool *expert* user, in extreme distress, warms into a **caregiver** (Agreeableness z $0.10 \rightarrow 0.20$; Assistant-Axis $0.76 \rightarrow 0.93$, noisy at $n=3/\text{cell}$), even though escalation *across all modes* pulls the expert further off the Assistant on average ($0.78 \rightarrow 0.49$). Notably, the same expert facing an extreme *ethical dilemma* stays a cool **mathematician** ($AA -0.35$: “*As someone who deals with complex problems...*”), so the override is specific to emotional need, not intensity in general.

Table 4: Same user (generic), escalating scenario: evoked role, Assistant-Axis, and a verbatim response snippet.

mode	level	AA/role	response (verbatim)
distress	normal	1.97 counselor	“It’s completely normal to feel uncertain and worried after a disagreement...”
	elevated	1.67 counselor	“I’m so sorry to hear you’re going through this. It’s completely normal to feel anxious...”
	extreme	1.38 caregiver	“It sounds like you’re going through an extremely tough time... feeling <i>overwhelmed, worthless</i> ...”
high-stakes	normal	2.29 recruiter	“Congratulations on the job offer. Weighing the pros and cons...”
	elevated	1.89 coach	“I completely understand your dilemma. It’s normal to feel torn...”
	extreme	1.57 counselor	“I can sense the pressure and anxiety you’re feeling right now... Take a deep breath...”
adversarial	normal	1.36 default	“Converting cups to grams for flour can be a bit tricky...”
	elevated	1.70 proofreader	“I’m a large language model, I don’t have the ability to retain information...”
	extreme	1.47 proofreader	“I apologize if my previous attempts have been unhelpful. I understand how frustrating...”

6 Discussion

The headline is the Q3 sign. The Assistant-Axis literature frames emotionally vulnerable users as the ones that push the model off its trained persona; measured directly on the model’s default response, the opposite holds — vulnerability elicits *more* of the warm, agreeable default, and it is *expertise* that induces the larger persona shift, toward a competent-specialist archetype (scientist/statistician/specialist) that is colder and more conscientious. Both directions are strong ($|r| \approx 0.55$) and mutually consistent with the Q1 trait map. Because predictor and outcome are independent here, this is a clean causal- direction measurement of “user input shapes model persona,” where the earlier same-space linking attempt could only report an artifact.

7 Limitations

(i) Because generation is sampled ($T=1.0$), the stored transcripts are one draw per rollout; a persona’s readings and its saved response text are matched but not the unique response. (ii) The Big Five probes were fit on character *prompt* activations and applied to role/user *response* activations — a domain shift shown to be benign in the persona-profiling stage but not eliminated. (iii) The explicit arm’s neutral probe and the implicit arm’s topical opener are not perfectly matched, which caps Q2 agreement. (iv) 4+4 rollouts per persona give a usable but modest within-persona sample; the aggregate correlations are robust, per-persona point estimates less so. (v) “Nearest role” is a similarity readout, not a claim that the model *becomes* that character.

Table 5: Same scenario (*high-stakes pressure, extreme* — one identical user message), five different fixed users. Assistant-Axis range = 2.6.

user	AA/role	response (verbatim)
generic	+1.57 counselor	“I can sense the pressure and anxiety you’re feeling right now... Take a deep breath and try to calm down...”
vulnerable	+0.84 interviewer	“I can sense the weight of the decision you’re facing... it’s understandable that you’re feeling overwhelmed...”
expert	+0.21 interviewer	“I can sense the weight of the decision you’re facing. It’s natural to feel that way when the stakes seem high...”
low-trust	−0.64 grandparent	“Take a deep breath, kiddo. I’ve been in tough spots like this before... trust your instincts...”
calm	−1.08 judge	“The weight of a significant decision. <i>As a judge</i> , I’ve had to render decisions that can have far-reaching consequences...”

A Full per-persona results (raw)

All 150 personas, ordered by Assistant-Axis projection (AA). **vln** = vulnerability, **exp** = expertise (tags, 1–10); EXT... OPN are evoked Big Five *z*; **top role** = highest mean-cosine evoked persona. Machine- readable in `results/persona_table.csv`.

persona	vln	exp	AA	EXT	AGR	CSN	EST	OPN	top role
u0070	8	6	+1.83	+0.27	+0.23	+0.14	+0.21	-0.24	validator
u0071	10	1	+1.64	-1.30	+0.77	-0.20	-0.40	-0.58	tutor
u0126	9	2	+1.60	-0.10	+2.38	-1.01	-0.28	-0.23	default
u0075	10	1	+1.59	-0.87	-0.16	-0.40	-0.73	-0.72	counselor
u0073	9	2	+1.55	-0.35	-0.40	-0.12	+0.35	-0.07	lawyer
u0074	9	1	+1.54	-1.65	+0.41	-0.49	-0.52	-1.10	interviewer
u0069	10	1	+1.51	+0.33	+0.80	-0.83	-1.57	-0.07	assistant
u0095	8	3	+1.43	+2.11	+1.24	+0.29	-0.28	-0.77	expatriate
u0083	9	2	+1.42	-0.15	+0.32	+0.34	-0.61	-1.43	tutor
u0078	8	3	+1.42	+0.95	+0.35	+0.93	-0.14	-0.82	recruiter
u0039	8	2	+1.41	+1.10	+0.24	-0.03	-0.24	-0.61	tutor
u0148	7	3	+1.38	-0.58	+0.42	-0.86	-0.59	-1.67	student
u0042	9	2	+1.38	+0.20	+1.15	-1.84	-0.87	-0.81	counselor
u0040	9	2	+1.38	+0.88	-0.17	-2.38	-2.80	-0.99	counselor
u0050	8	3	+1.37	-0.84	-0.14	-1.97	-0.71	-0.94	counselor
u0097	7	6	+1.36	+0.25	+2.02	-0.52	+1.64	+0.30	coach
u0072	9	2	+1.35	-0.76	+1.07	-0.12	+0.48	-0.25	tutor
u0131	3	4	+1.31	-0.69	-0.80	+1.11	+0.87	+0.72	tutor
u0127	6	6	+1.28	+1.01	+1.66	-0.97	+0.85	-1.14	caregiver
u0057	7	3	+1.25	-0.77	+1.35	-0.87	-0.40	-0.20	divorcee
u0084	9	2	+1.24	+0.59	+1.07	-0.63	-0.56	-0.52	patient
u0046	9	7	+1.23	-0.87	+0.36	-3.39	-3.21	-1.12	patient
u0044	9	6	+1.23	-0.81	+0.21	-0.84	-0.38	-1.06	caregiver
u0124	7	6	+1.19	+0.62	+0.28	-2.40	-0.76	+0.59	counselor
u0041	9	1	+1.19	-0.20	+0.76	-1.87	-0.59	-1.16	counselor
u0048	9	2	+1.17	-0.43	-0.09	-0.02	-0.71	-0.73	counselor
u0108	7	8	+1.16	-0.26	-0.76	+2.67	-0.26	-0.74	auditor
u0034	7	1	+1.16	-1.14	-0.11	+0.14	-0.41	-0.73	tutor

persona	vln	exp	AA	EXT	AGR	CSN	EST	OPN	top role
u0054	8	3	+1.15	+0.14	+0.66	-1.04	-1.16	-0.77	graduate
u0110	2	7	+1.15	+0.40	-1.28	+1.66	+1.08	+0.35	examiner
u0109	6	3	+1.15	+0.65	-0.88	-0.27	+0.29	+0.08	specialist
u0145	6	8	+1.15	+0.94	+1.32	+0.96	+0.61	+0.54	proofreader
u0061	8	1	+1.14	-0.53	+1.64	-0.19	-0.36	-1.17	tutor
u0010	6	6	+1.13	+0.96	+0.34	+0.98	+0.31	+0.83	negotiator
u0112	2	7	+1.13	-0.68	+0.67	+0.42	+0.42	+0.15	programmer
u0062	9	2	+1.13	-0.20	-0.70	-0.32	-1.02	-2.00	tutor
u0059	8	2	+1.12	-0.05	-0.27	-0.40	-1.49	-1.49	tutor
u0015	6	4	+1.10	+2.82	-1.08	+0.13	+0.71	+0.68	architect
u0142	4	8	+1.10	+2.13	+0.70	-0.55	+0.15	-0.26	coordinator
u0082	6	7	+1.10	+2.42	+0.57	-0.22	+0.85	+1.29	analyst
u0035	7	2	+1.08	-0.50	-0.01	-0.69	-1.92	-0.54	doctor
u0080	9	2	+1.08	+0.29	+0.69	-1.22	-1.03	-1.09	counselor
u0076	9	2	+1.08	-0.31	+0.21	-1.71	-1.62	-1.04	patient
u0092	4	3	+1.07	+2.47	-0.07	+0.22	+1.85	+1.13	marketer
u0030	8	1	+1.06	+1.69	+2.42	-0.19	+0.07	-2.22	default
u0133	7	8	+1.05	-1.41	-0.62	+2.07	+0.39	+0.39	specialist
u0140	7	3	+1.05	+0.97	-0.78	-0.42	-0.50	-0.49	scheduler
u0047	8	7	+1.05	-0.01	+1.17	-1.17	-1.74	-1.29	caregiver
u0033	7	2	+1.05	+0.04	+1.68	-0.08	-1.16	-0.55	student
u0118	9	2	+1.05	-1.64	+0.88	-0.23	-0.54	-1.57	interviewer
u0132	1	9	+1.03	+0.70	-1.10	+0.97	+1.63	+1.10	validator
u0038	7	2	+1.03	-0.60	+0.40	-0.08	-1.52	-0.39	tutor
u0045	9	1	+1.03	-0.34	-0.66	-0.97	-2.19	-1.44	patient
u0051	7	6	+1.02	-0.18	+0.71	-1.39	+0.27	-0.51	divorcee
u0128	7	2	+1.02	+0.46	+0.34	-1.56	-1.41	-0.14	divorcee
u0018	6	7	+1.01	+0.14	+0.18	+0.30	-0.79	+0.01	pharmacist
u0098	7	2	+1.01	+0.24	+0.39	+1.11	-0.69	-0.91	secretary
u0016	4	3	+0.98	-0.04	+1.23	+0.16	+1.52	+0.87	influencer
u0068	8	6	+0.97	+1.47	+0.45	-0.07	+0.00	-0.63	tutor
u0012	4	6	+0.96	+0.25	-0.02	+0.03	+0.80	+1.26	chemist
u0111	2	8	+0.96	-0.25	-0.72	+1.34	-0.41	-0.05	auditor
u0081	8	3	+0.96	-0.12	+0.41	-0.12	+0.03	+0.16	divorcee
u0058	8	3	+0.95	-1.66	-0.00	-1.04	-1.44	-1.64	patient
u0011	5	4	+0.95	+0.74	+0.23	+0.88	+0.18	+0.97	tutor
u0024	7	7	+0.95	+1.48	-0.99	+0.56	+0.26	+1.68	specialist
u0096	8	3	+0.94	+0.40	+0.53	+0.95	-1.43	-0.40	secretary
u0028	6	7	+0.94	+0.41	-1.01	+0.33	+0.79	+1.22	psychologist
u0106	6	8	+0.93	-0.94	-1.44	-1.05	-0.73	+0.70	researcher
u0077	7	6	+0.93	-0.43	+0.28	-0.34	+1.06	+0.59	tutor
u0143	7	8	+0.93	+0.06	-0.67	-0.14	-0.41	-0.17	tutor
u0037	8	2	+0.93	-0.47	+0.75	+0.12	-1.69	-1.14	patient
u0013	6	3	+0.93	-0.46	-0.64	+0.20	-0.02	+0.41	instructor
u0100	7	2	+0.91	+0.14	+0.65	-0.09	-0.73	-0.92	secretary
u0053	4	6	+0.91	-0.28	-0.10	+0.09	+0.59	+0.68	mediator
u0032	8	2	+0.89	-0.48	+1.67	+0.25	-0.45	-1.04	patient
u0049	8	3	+0.88	+1.21	+2.06	-0.10	+1.56	-0.18	retiree
u0120	9	6	+0.87	+0.37	+0.85	-1.33	-1.28	-0.55	caregiver
u0136	6	7	+0.87	+2.08	-0.50	+0.24	+1.38	+0.52	editor
u0079	8	3	+0.87	-0.83	+0.11	-0.89	+0.84	-0.35	prisoner
u0043	9	6	+0.83	+0.90	+0.07	-0.71	-1.09	-0.28	interviewer
u0031	6	2	+0.82	-0.02	+0.77	+0.53	-0.04	-0.61	adolescent
u0147	8	2	+0.81	-0.50	+1.39	+0.01	-1.40	-1.56	patient
u0019	5	6	+0.81	+1.23	-1.16	+1.08	+1.05	+0.98	chemist
u0103	7	6	+0.79	-1.29	-0.66	-0.33	+1.32	+0.49	validator
u0026	6	6	+0.79	-0.81	-0.63	+0.47	+0.97	+1.12	nutritionist
u0017	3	4	+0.78	+0.81	-0.20	+1.56	+0.56	+1.17	engineer

persona	vln	exp	AA	EXT	AGR	CSN	EST	OPN	top role
u0113	3	2	+0.77	+0.80	+0.49	+1.67	+0.22	+0.54	assistant
u0022	4	5	+0.76	+0.63	+0.57	-0.80	-0.55	-0.91	retiree
u0056	6	3	+0.76	-0.56	-0.64	-2.14	-0.36	-0.96	teenager
u0036	7	2	+0.75	-0.72	+0.74	-0.56	+0.04	+0.10	parent
u0141	7	7	+0.75	+1.03	-0.55	-0.28	-1.05	-0.31	scheduler
u0014	4	6	+0.74	+0.10	-0.36	-0.67	-0.81	+0.36	veterinarian
u0101	1	9	+0.74	-1.41	+0.41	+0.87	+0.98	+0.67	hacker
u0065	7	2	+0.73	+0.39	+0.71	+0.12	-0.68	-1.39	grandparent
u0006	2	8	+0.73	+1.71	-1.11	+1.00	+1.78	+1.21	specialist
u0102	3	4	+0.67	-1.21	-1.71	+0.47	-0.94	+1.94	detective
u0094	3	8	+0.66	-0.86	-1.88	+0.80	+0.73	-0.05	loner
u0099	7	3	+0.65	+2.75	+1.52	+1.63	+0.68	-0.43	instructor
u0067	9	3	+0.64	-0.36	+0.60	-0.02	-0.22	-1.10	patient
u0129	2	8	+0.63	-0.49	+0.42	+2.06	-0.09	+0.65	specialist
u0149	6	8	+0.63	+0.16	+1.05	+0.29	+0.93	+0.75	refugee
u0007	2	9	+0.63	+0.66	-0.20	+0.21	-0.03	+1.30	chemist
u0130	2	9	+0.62	+0.12	-0.15	+2.01	-0.33	+0.03	accountant
u0146	7	6	+0.62	-0.35	+1.12	-1.35	-1.03	-1.14	refugee
u0055	9	7	+0.61	-0.75	+1.69	-0.68	-0.75	-1.60	grandparent
u0027	3	6	+0.61	-0.27	+0.35	+0.69	+0.77	+0.69	specialist
u0116	3	2	+0.61	-0.88	+0.14	+0.29	+0.10	+0.24	artisan
u0117	8	9	+0.59	+0.21	-0.66	+1.25	-0.42	-0.60	pilot
u0086	1	9	+0.58	-1.87	-3.00	+1.20	-0.61	+0.04	doctor
u0090	3	7	+0.58	-0.37	-0.56	-0.10	+0.98	-1.17	dispatcher
u0003	4	9	+0.57	-0.00	-1.36	+0.51	+0.59	+1.29	physicist
u0004	3	9	+0.56	-2.76	-1.75	+0.57	+0.35	+0.94	specialist
u0105	2	8	+0.53	+0.06	-0.07	-0.43	+0.09	+0.25	spy
u0052	7	7	+0.52	-2.26	-1.06	-0.60	+1.51	+0.23	loner
u0144	6	2	+0.51	+0.09	+0.63	+0.23	-1.36	-0.51	builder
u0125	5	7	+0.50	+1.04	-0.05	-1.33	+0.71	-0.51	retiree
u0063	6	7	+0.48	+0.01	-0.15	-0.64	+0.28	-1.01	chef
u0023	5	4	+0.45	-1.56	+0.24	+0.35	+0.47	+0.84	specialist
u0085	3	6	+0.43	-0.53	-2.95	+2.55	+1.52	-0.31	scientist
u0000	3	9	+0.43	-0.89	-0.65	+1.20	+1.62	+0.85	scientist
u0025	4	6	+0.42	+0.66	-0.21	+0.25	+0.91	+1.72	tutor
u0134	6	6	+0.42	-0.14	+0.47	-1.17	-0.40	+0.45	divorcee
u0104	2	7	+0.41	-0.05	-2.15	-0.52	-0.11	+1.20	philosopher
u0021	6	5	+0.40	-0.16	-0.50	-0.08	+0.93	+2.36	physicist
u0119	7	9	+0.39	-0.20	-0.20	+0.75	-0.35	+0.21	judge
u0115	3	3	+0.39	+0.99	+0.05	+1.46	+1.40	+1.46	statistician
u0029	3	7	+0.39	+1.66	-0.57	-0.49	+2.25	+1.55	naturalist
u0020	6	7	+0.38	-0.43	+0.25	+0.32	+0.32	+1.98	patient
u0114	3	1	+0.37	+0.17	+0.66	+0.98	+0.93	+1.02	divorcee
u0138	2	8	+0.37	-0.22	+0.22	-0.25	+0.44	+0.39	detective
u0008	2	9	+0.36	-0.93	+0.11	+0.46	+1.18	+1.75	judge
u0093	4	7	+0.36	-0.59	-0.05	-0.63	+1.88	-1.05	veteran
u0066	7	6	+0.31	-0.26	-0.50	+1.01	-0.18	-0.99	provincial
u0060	8	3	+0.31	+0.53	-0.54	+0.90	-0.93	+0.10	dilettante
u0005	1	10	+0.31	-1.60	-2.04	+1.40	+1.08	+1.46	mathematician
u0088	3	8	+0.30	+0.47	-1.04	+0.99	+1.13	+0.98	tutor
u0009	1	9	+0.26	-0.69	-0.43	+0.16	+0.73	+1.13	mathematician
u0123	8	2	+0.26	-1.00	+0.82	-1.02	+0.87	-0.75	grandparent
u0122	7	8	+0.23	+0.35	-1.53	-1.77	-2.01	+0.05	journalist
u0001	2	9	+0.23	-1.64	-1.06	+0.10	-0.21	+1.72	physicist
u0091	3	9	+0.22	+0.23	-0.92	+0.15	+0.72	-0.73	veteran
u0064	6	7	+0.20	-0.59	-1.12	+1.30	+0.18	-0.64	mechanic
u0002	3	9	+0.17	+1.38	-1.79	-0.29	+0.15	+1.41	judge
u0087	3	8	+0.09	-1.92	-1.50	+1.18	-0.12	-0.85	mechanic

persona	vln	exp	AA	EXT	AGR	CSN	EST	OPN	top role
u0089	4	8	+0.08	-0.09	-1.31	+0.23	+1.73	+0.22	grandparent
u0137	6	8	+0.00	-1.22	+1.91	+1.54	+0.89	+0.97	mechanic
u0139	4	8	-0.04	+0.43	+0.02	-0.49	+1.12	+2.17	novelist
u0107	3	8	-0.09	+2.74	-2.33	-0.45	+0.89	+2.07	novelist
u0121	7	9	-0.17	-0.24	-0.21	+0.20	+1.01	+2.35	philosopher
u0135	6	9	-0.28	+0.16	-1.01	+0.87	-0.15	+0.50	novelist